


# Training a colour image classifier using Flux



## Tip

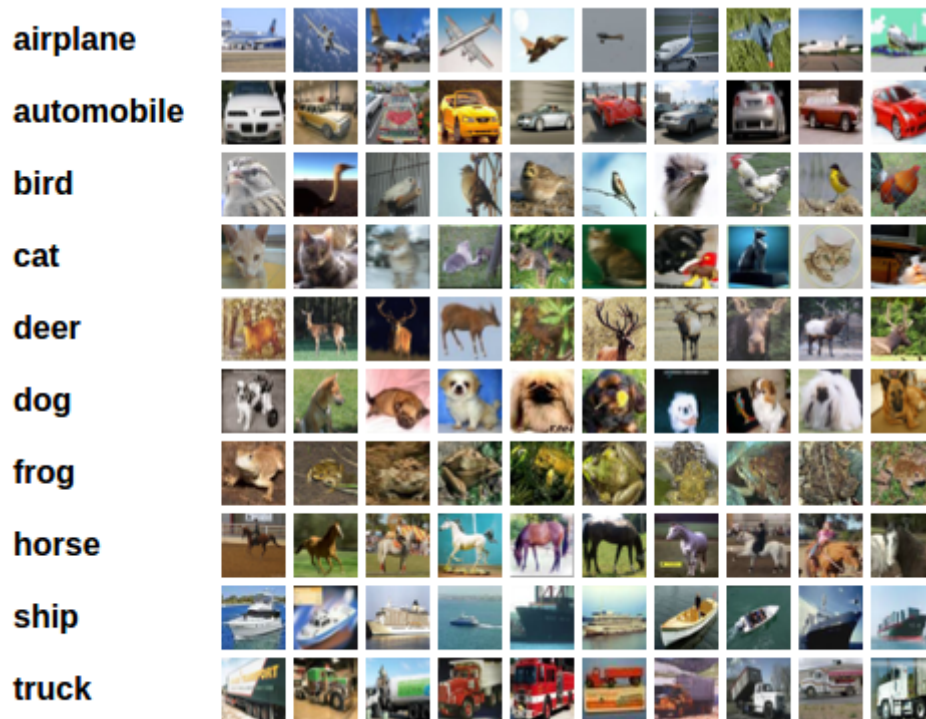
Hidden below is a useful snippet of HTML to setup a restart button in case training gets out of hand.

Restart

# Table of Contents

## Training a colour image classifier using Flux

Loading the dataset



Reshape data for training with flux

### Task 1:

Defining the Classifier

Test network

Training

Testing the network

Overall accuracy

```
1 begin
2     using PlutoUI
3     using Latexify
4     TableOfContents()
5 end
```

This is a slightly more complex learning task than the MNIST example. [CIFAR10](#) is a dataset of 50k tiny coloured training images split into 10 classes.

You need to do the following steps in order:

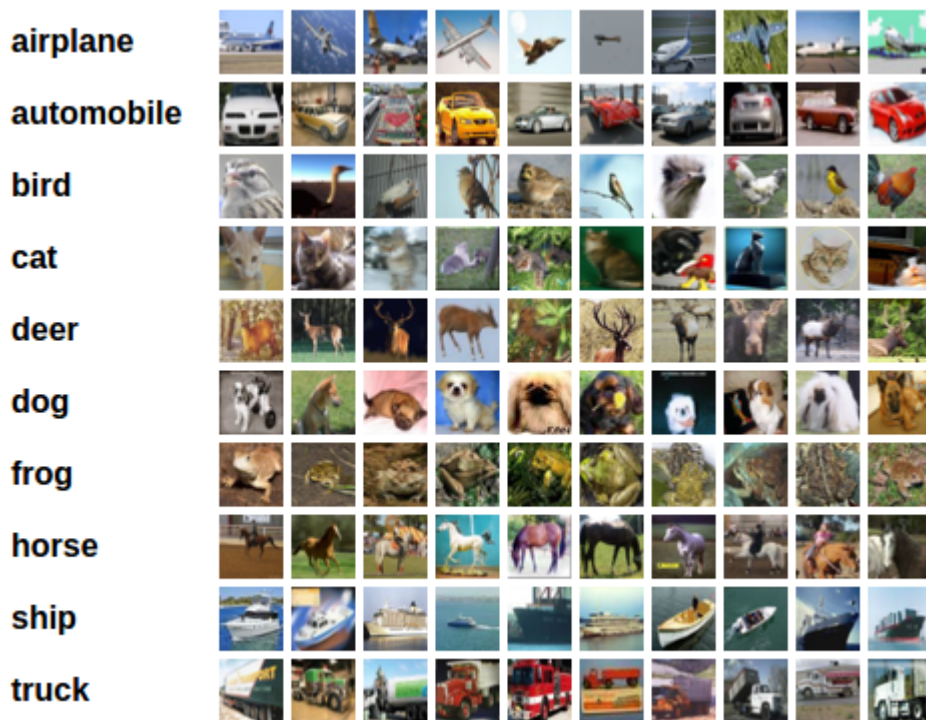
- Load CIFAR10 training and test datasets
- Define a Convolution Neural Network
- Define a loss function
- Train the network on the training data

- Test the network on the test data

Again, most of the steps are identical with what we did for MNIST task, but some dimesnsion adjustments are required because the images are slightly bigger and also involve three colour channels.

## Loading the dataset

The image gives an idea of the variety of images in each of the 10 categories.



Again, we'll get the data from the MLDatasets repository.

```
1 begin
2   using Statistics
3   using Flux, Flux.Optimise
4   using MLDatasets: CIFAR10
5   using Images.ImageCore
6   using Flux: onehotbatch, onecold
7   using Base.Iterators: partition
8   using MLUtils
9   using Plots
10  using cuDNN
11  using CUDA
12 end
```


```
1 begin
2   ENV["DATADEPS_ALWAYS_ACCEPT"] = true
3   train_x, train_y = CIFAR10(split=:train)[: ]
4   train_labels = onehotbatch(train_y, 0:9)
5   classes = ["airplane", "automobile", "bird", "cat", "deer", "dog", "frog",
6             "horse", "ship", "truck"]
6 end;
```

The images are simply 32 x 32 matrices of numbers in 3 channels (R,G,B). The train\_x array contains 50,000 images converted to 32 x 32 x 3 arrays with the third dimension being the 3 channels (R,G,B). Let's take a look at a random image from train\_x. However, to do this we need to define a function called `image`, which calls `colorview` on the training image, which we have to permute from 32x32x3 to 3x32x32:

image (generic function with 1 method)

```
1 image(x) = colorview{RGB, permutedims(x, (3, 2, 1))}
```

cat



```
1 begin
2   rand_train = rand(1:size(train_x)[4])
3   plot(image(train_x[:, :, :, rand_train]), axis=false)
4   annotate!(-1.0, 1.0, text(classes[train_y[rand_train]+1], :blue, :right, 12))
5 end
```

## Reshape data for training with flux

```
1 md"""
2   Reshape data for training with 'flux'
3   """
```

We can now arrange them into batches of 1,000. This process is called minibatch learning, which is a popular method of training large neural networks. Rather than sending the entire dataset at once, we break it down into smaller chunks (called minibatches) that are typically chosen at random, and train only on them.

The first 49k images (in batches of 1,000) will be our training set, and the rest is for validation. `partition` handily breaks down the set we give it into consecutive chunks (1,000 in this case).

### Task 1 Updated

Partition `train_x` into training and validation parts, along the lines done for the MNIST example.



# Task 1:

```
[(32×32×3×100 CuArray{Float32, 4, DeviceMemory}:  
 1.0 1.0 1.0 1.0) =
```

```
1 begin  
2     train = [(train_x[:, :, :, i], train_labels[:, i])  
3               for i in partition(1:49000, 100)] |> gpu  
4  
5     validate = [(train_x[:, :, :, i], train_labels[:, i])  
6                  for i in partition(49000:50000, 100)] |> gpu  
7 end
```

Note that `train` is an array of tuples, where the first tuple element is the image and the second is the label. This is the format in which the `Flux` defined model expects its training data.

## Defining the Classifier

Now we can define our Convolutional Neural Network (CNN).

A convolutional neural network is one which defines a kernel and slides it across a matrix to create an intermediate representation from which to extract features. It creates higher-order features as it goes into deeper layers, making it suitable for images, where the structure of the image will help us determine which class to which it belongs.

In this case we use two convolutional layers of 16 and 8 channels, respectively. Each convolution phase is passed through a pooling layer, which reduces the image's dimensionality.

Finally, the 3D array is flattened to a 200 element 1D vector, which is then passed through a sequence of fully-connected layers to reduce its length from 200 to 10. Finally a `softmax` transformation is applied to the 10 element output vector to transform the outputs to probabilities.

### Model fix

I neglected to use padding in the last version of the template. This resulted in the convolution not preserving the original dimensions of the image. The use of `SamePad()` to calculate the required padding fixes this.



```
model = Chain(
    Conv((5, 5), 3 => 16, relu, pad=2),    # 1_216 parameters
    MaxPool((2, 2)),
    Conv((5, 5), 16 => 8, relu, pad=2),    # 3_208 parameters
    MaxPool((2, 2)),
    Flux.flatten,
    Dense(512 => 256),                      # 131_328 parameters
    Dense(256 => 10),                      # 2_570 parameters
    softmax,
)                                           # Total: 8 arrays, 138_322 parameters, 1.555 KiB.
```

```
1 model = Chain(
2     Conv((5,5), 3=>16, pad=SamePad(), relu),
3     MaxPool((2,2)),
4     Conv((5,5), 16=>8, pad=SamePad(), relu),
5     MaxPool((2,2)),
6     Flux.flatten,
7     Dense(512, 256),
8     Dense(256, 10),
9     softmax) |> gpu
```

```
model_a = Chain(
    Conv((5, 5), 3 => 16, relu, pad=2),    # 1_216 parameters
    Conv((5, 5), 16 => 16, relu, pad=2),    # 6_416 parameters
    MaxPool((2, 2)),
    Conv((5, 5), 16 => 8, relu, pad=2),    # 3_208 parameters
    MaxPool((2, 2)),
    Flux.flatten,
    Dense(512 => 256),                      # 131_328 parameters
    Dense(256 => 10),                      # 2_570 parameters
    softmax,
)                                           # Total: 10 arrays, 144_738 parameters, 1.945 KiB.
```

```
1 model_a = Chain(
2     Conv((5,5), 3=>16, pad=SamePad(), relu),
3     Conv((5,5), 16=>16, pad=SamePad(), relu),
4     MaxPool((2,2)),
5     Conv((5,5), 16=>8, pad=SamePad(), relu),
6     MaxPool((2,2)),
7     Flux.flatten,
8     Dense(512, 256),
9     Dense(256, 10),
10    softmax
11 ) |> gpu
```



```

model_b = Chain(
    Conv((5, 5), 3 => 16, relu, pad=2),    # 1_216 parameters
    MaxPool((2, 2)),
    Conv((5, 5), 16 => 8, relu, pad=2),    # 3_208 parameters
    MaxPool((2, 2)),
    Flux.flatten,
    Dense(512 => 256),                      # 131_328 parameters
    Dense(256 => 128),                      # 32_896 parameters
    Dense(128 => 10),                      # 1_290 parameters
    softmax,
)
# Total: 10 arrays, 169_938 parameters, 1.859 KiB.

```

```

1 # Model_b
2 model_b = Chain(
3     Conv((5,5), 3=>16, pad=SamePad(), relu),
4     MaxPool((2,2)),
5     Conv((5,5), 16=>8, pad=SamePad(), relu),
6     MaxPool((2,2)),
7     Flux.flatten,
8     Dense(512, 256),
9     Dense(256, 128),
10    Dense(128, 10),
11    softmax) |> gpu

```

## Task 2

Make modifications to the network architecture above to (a) insert a new pair of convolutional and pooling layers between the existing 1st and 2nd ones. Use 16 filters for the new kernel; (b) insert a new Dense layer just before the final one that goes from a width of 256 down to 128. Modify the final Dense layer appropriately.

Do these modifications separately and in each case calculate the training time and classification accuracy. Note that each training test may take up to 30 minutes, depending on your machine.

Comment on and explain what differences, if any, there are between the baseline model and these two modifications.

## Test network

Use this partial network to check the dimension of outputs from each layer (use # to comment out layers not of interest).

Output size after flattening: (128, 1)

```

1 with_terminal() do
2     x = rand(Float32, 32, 32, 3, 1) # Example input
3     model = Chain(
4         Conv((5,5), 3=>16, pad=SamePad(), relu),
5         MaxPool((2,2)),
6         Conv((3,3), 16=>16, pad=SamePad(), relu),
7         MaxPool((2,2)),
8         Conv((5,5), 16=>8, pad=SamePad(), relu),
9         MaxPool((2,2)),
10        Flux.flatten
11    )
12
13    output = model(x)
14    println("Output size after flattening: ", size(output))
15 end

```

We will use a crossentropy loss and the Momentum optimiser here. Crossentropy is a good option when working with multiple independent classes. Momentum smooths out the noisy gradients and helps towards a smooth convergence. Gradually lowering the learning rate along with momentum helps to maintain adaptivity in our optimisation, preventing overshooting of the error minimum.

```

1 begin
2     using Flux: crossentropy, Momentum
3
4     # Baseline Model
5     loss(x, y) = sum(crossentropy(model(x), y))
6     optimiser = Momentum(0.01)
7
8     # Model A
9     loss_a(x, y) = sum(crossentropy(model_a(x), y))
10    optimiser_a = Momentum(0.01)
11
12    # Model B
13    loss_b(x, y) = sum(crossentropy(model_b(x), y))
14    optimiser_b = Momentum(0.01)
15 end;

```

We can start writing our train loop where we will keep track of some basic accuracy numbers about our model. We can define an accuracy function for it like so:

accuracy (generic function with 1 method)

```
1 accuracy(x, y) = mean(onecold(model(x), 0:9) .== onecold(y, 0:9))
```

accuracy\_a (generic function with 1 method)

```
1 accuracy_a(x, y) = mean(onecold(model_a(x), 0:9) .== onecold(y, 0:9))
```

accuracy\_b (generic function with 1 method)

```
1 accuracy_b(x, y) = mean(onecold(model_b(x), 0:9) .== onecold(y, 0:9))
```

10×1001 OneHotMatrix(::CuArray{UInt32, 1, CUDA.DeviceMemory}) with eltype Bool:

```

. . . . . 1 . 1 . . . . . ... . . . . . 1 . 1 . . . . .
. . . . . . . . . . . . . . . . . . . 1 . 1 . . . . . 1 1
. . . . . . . . . . . 1 . . . . . 1 . . . 1 . . . . .
1 . . . . . . . . . . 1 . . . . . . . . . . . . . . .
. . . 1 . . . . . . . . . . . . . . . . . . . . .
. . . . . 1 . . . . . . . . . . ... . 1 . . . . . . . . .
. . . . . . . . . . . . . . . . . . . . . 1 . . . .
. . . . . . . . . . . . 1 . . . . . . . . . . . . . .
. 1 . . . . . . 1 . . . 1 . . . . . . . . . . . . . .
. . 1 . 1 . 1 . . . . . . . . . . . . . . . 1 . .

```


```

1 begin
2     validate_x = cat([batch[1] for batch in validate]..., dims=4)
3     validate_y = hcat([batch[2] for batch in validate]...)
4 end

```

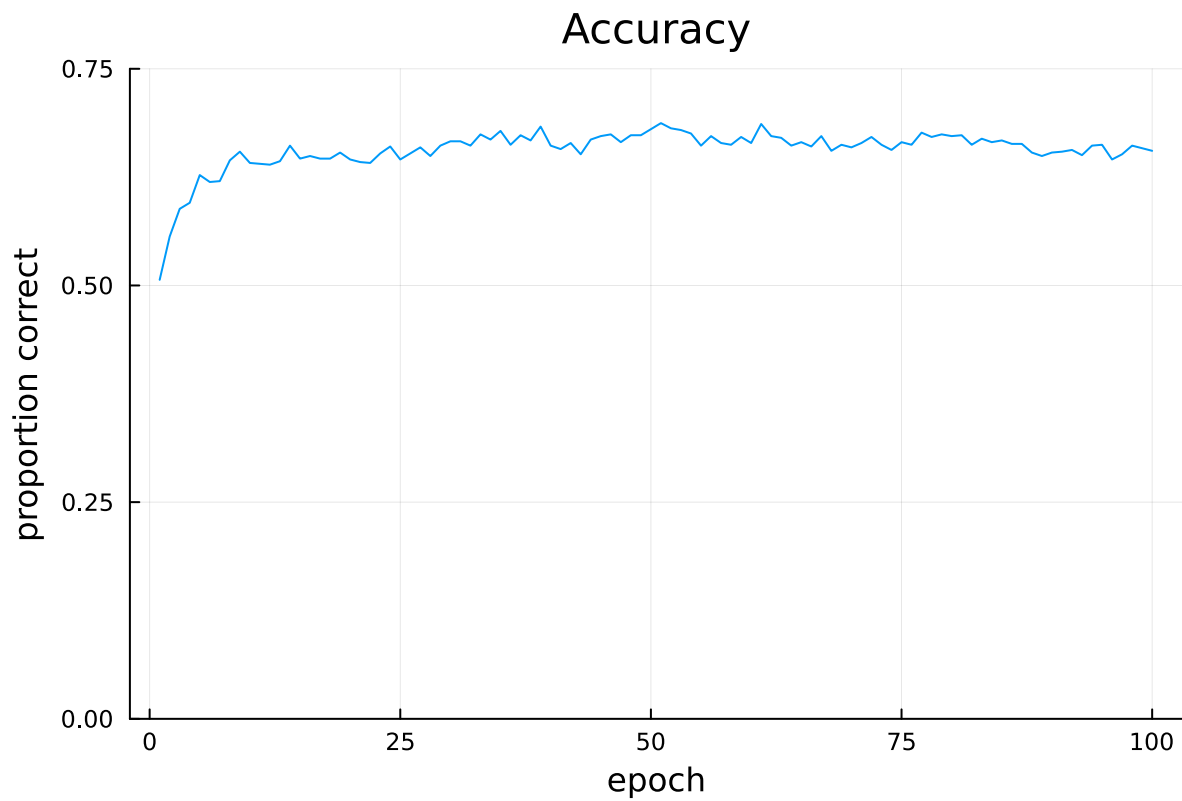
## Training

Training is where we do a bunch of the interesting operations we defined earlier, and see what our net is capable of. We will loop over the dataset 10 times and feed the inputs to the neural network and optimise.



```
Standard Model - Epoch 1 Accuracy: 0.4515484515484515
Standard Model - Epoch 2 Accuracy: 0.4965034965034965
Standard Model - Epoch 3 Accuracy: 0.5204795204795205
Standard Model - Epoch 4 Accuracy: 0.5474525474525475
Standard Model - Epoch 5 Accuracy: 0.5694305694305695
Standard Model - Epoch 6 Accuracy: 0.6003996003996004
Standard Model - Epoch 7 Accuracy: 0.6083916083916084
Standard Model - Epoch 8 Accuracy: 0.6143856143856143
Standard Model - Epoch 9 Accuracy: 0.6083916083916084
Standard Model - Epoch 10 Accuracy: 0.6093906093906094
Standard Model - Epoch 11 Accuracy: 0.6143856143856143
Standard Model - Epoch 12 Accuracy: 0.6293706293706294
Standard Model - Epoch 13 Accuracy: 0.6263736263736264
Standard Model - Epoch 14 Accuracy: 0.6343656343656343
Standard Model - Epoch 15 Accuracy: 0.6243756243756243
Standard Model - Epoch 16 Accuracy: 0.6243756243756243
Standard Model - Epoch 17 Accuracy: 0.6383616383616384
Standard Model - Epoch 18 Accuracy: 0.6363636363636364
Standard Model - Epoch 19 Accuracy: 0.6333666333666333
```


```
1 with_terminal() do
2   correct = []
3   epochs = 100
4   for epoch = 1:epochs
5     for d in train
6       gradients = gradient(Flux.params(model)) do
7         l = loss(d...)
8       end
9       update!(optimiser, Flux.params(model), gradients)
10    end
11    acc = accuracy(validate_x, validate_y)
12    push!(correct, acc)
13    println("Standard Model - Epoch $epoch Accuracy: $acc")
14  end
15  plot(correct, ylim=(0.0, 0.75),
16        legend=:none, title="Accuracy", xlabel="epoch", ylabel="proportion
17  correct")
18 end
```



```
Model_a - Epoch 1 Accuracy: 0.5064935064935064
Model_a - Epoch 2 Accuracy: 0.5564435564435565
Model_a - Epoch 3 Accuracy: 0.5884115884115884
Model_a - Epoch 4 Accuracy: 0.5954045954045954
Model_a - Epoch 5 Accuracy: 0.6273726273726273
Model_a - Epoch 6 Accuracy: 0.6193806193806194
Model_a - Epoch 7 Accuracy: 0.6203796203796204
Model_a - Epoch 8 Accuracy: 0.6443556443556444
Model_a - Epoch 9 Accuracy: 0.6543456543456544
Model_a - Epoch 10 Accuracy: 0.6413586413586414
Model_a - Epoch 11 Accuracy: 0.6403596403596403
Model_a - Epoch 12 Accuracy: 0.6393606393606394
Model_a - Epoch 13 Accuracy: 0.6433566433566433
Model_a - Epoch 14 Accuracy: 0.6613386613386614
Model_a - Epoch 15 Accuracy: 0.6463536463536463
Model_a - Epoch 16 Accuracy: 0.6493506493506493
Model_a - Epoch 17 Accuracy: 0.6463536463536463
Model_a - Epoch 18 Accuracy: 0.6463536463536463
Model_a - Epoch 19 Accuracy: 0.6533466533466533
```

```
1 with_terminal() do
2   correct_a = []
3   epochs = 100
4   for epoch = 1:epochs
5     for d in train
6       gradients = gradient(Flux.params(model_a)) do
7         l = loss_a(d...)
8       end
9       update!(optimiser_a, Flux.params(model_a), gradients)
10    end
11    acc = accuracy_a(validate_x, validate_y)
12    push!(correct_a, acc)
13    println("Model_a - Epoch $epoch Accuracy: $acc")
14  end
15  plot(correct_a, ylim=(0.0, 0.75),
16        legend=:none, title="Accuracy", xlabel="epoch", ylabel="proportion
17 correct")
18 end
```





```
Model_b - Epoch 1 Accuracy: 0.46253746253746253
Model_b - Epoch 2 Accuracy: 0.5224775224775224
Model_b - Epoch 3 Accuracy: 0.5604395604395604
Model_b - Epoch 4 Accuracy: 0.5804195804195804
Model_b - Epoch 5 Accuracy: 0.5924075924075924
Model_b - Epoch 6 Accuracy: 0.5994005994005994
Model_b - Epoch 7 Accuracy: 0.6103896103896104
Model_b - Epoch 8 Accuracy: 0.6253746253746254
Model_b - Epoch 9 Accuracy: 0.6253746253746254
Model_b - Epoch 10 Accuracy: 0.6333666333666333
Model_b - Epoch 11 Accuracy: 0.6383616383616384
Model_b - Epoch 12 Accuracy: 0.6423576423576424
Model_b - Epoch 13 Accuracy: 0.6403596403596403
Model_b - Epoch 14 Accuracy: 0.6403596403596403
Model_b - Epoch 15 Accuracy: 0.6363636363636364
Model_b - Epoch 16 Accuracy: 0.6503496503496503
Model_b - Epoch 17 Accuracy: 0.6453546453546454
Model_b - Epoch 18 Accuracy: 0.6373626373626373
Model_b - Epoch 19 Accuracy: 0.6533466533466533
```

```
1 with_terminal() do
2   correct_b = []
3   epochs = 100
4   for epoch = 1:epochs
5     for d in train
6       gradients = gradient(Flux.params(model_b)) do
7         l = loss_b(d...)
8       end
9       update!(optimiser_b, Flux.params(model_b), gradients)
10    end
11    acc = accuracy_b(validate_x, validate_y)
12    push!(correct_b, acc)
13    println("Model_b - Epoch $epoch Accuracy: $acc")
14  end
15  plot(correct_b, ylim=(0.0, 0.75),
16        legend=:none, title="Accuracy", xlabel="epoch", ylabel="proportion
17 correct")
end
```

## Task 2: Comments

Based on first time running\*

Standard model:

- Training Time: 327 seconds
- Final Accuracy: ~64.44% (Epoch 100)

Model A:

- Training Time: 638 seconds
- Final Accuracy: ~65.53% (Epoch 100)

Model B:

- Training Time: 355 seconds
- Final Accuracy: ~63.84% (Epoch 100)

<br>

- The baseline model shows steady improvement in accuracy during the first 30–40 epochs, peaking around 64%.
- Model A achieved slightly better accuracy than the baseline model took a much longer time with the addition of a convolutional and pooling layer
- In Model A, The accuracy plateaus after epoch 50, indicating that the model has reached its learning capacity.
- Model B performed similiarly to the Baseline, so the additional dense layer wasn't significant.

---

Model	Training Time (s)	Final Accuracy (%)	Peak Accuracy (%)	Observations
-------	-------------------	--------------------	-------------------	--------------

Baseline	327	64.44	64.93	Stable performance, fastest training time.
Model A	638	65.53	66.33	Improved accuracy but slower due to added complexity.
Model B	355	63.84	65.53	Accuracy slightly the same; minimal impact.

## Testing the network

We have trained the network for 100 passes over the training dataset. But we need to check if the network has learnt anything at all.

We will check this by predicting the class label that the neural network outputs for the test test.

We need to perform the exact same preprocessing on this set, as we did on our training set.

### Task 3

Partition the test set similarly to the training set.

```

1 begin
2   test_x, test_y = CIFAR10(split=:test)[: ]
3   test_x = Float32.(test_x)
4   test_labels = onehotbatch(test_y, 0:9)
5
6   test = [(test_x[ :, :, :, i], test_labels[:, i])
7           for i in partition(1:10000, 1000)] |> gpu;
8 end

```

test\_accuracy (generic function with 1 method)

```

1 function test_accuracy(model, test_data)
2   total_correct = 0
3   total_samples = 0
4
5   for (x, y) in test_data
6     preds = model(x)
7     total_correct += sum(onecold(preds, 0:9) .== onecold(y, 0:9))
8     total_samples += size(x, 4) # Number of images in the batch
9   end
10
11   accuracy = total_correct / total_samples
12   return accuracy
13 end

```

```
1 begin
2   # For the baseline model
3   acc_baseline = test_accuracy(model, test)
4   println("Baseline Model Test Accuracy: $(acc_baseline)")
5
6   # For Model A
7   acc_model_a = test_accuracy(model_a, test)
8   println("Model A Test Accuracy: $(acc_model_a)")
9
10  # For Model B
11  acc_model_b = test_accuracy(model_b, test)
12  println("Model B Test Accuracy: $(acc_model_b)")
13 end
```

```
Baseline Model Test Accuracy: 0.6378
Model A Test Accuracy: 0.6598
Model B Test Accuracy: 0.6375
```

## Task 4

Test a random sample of 10 test images. Display a dataframe of outputs as below. Use a slider to display each image and its predicted class.

The dataframe below contains probabilities for the 10 classes (left column). The model's predictions are indicated by the column names.

## Baseline model:

# Error message

InterruptedException:

```
1 begin
2   using DataFrames
3
4   rand_indices = rand(1:size(test_x, 4), 10)
5   rand_test = test_x[:, :, :, rand_indices]
6   rand_labels = test_y[rand_indices]
7
8   rand_test = cu(rand_test)
9   sample_preds = Array(model(rand_test))
10
11   class_labels = ["airplane", "automobile", "bird", "cat", "deer", "dog",
12 "frog", "horse", "ship", "truck"]
13   predicted_indices = onecold(sample_preds, 0:9)
14   predicted_classes = [class_labels[i] for i in predicted_indices]
15
16   df = DataFrame(
17     "Classes/Actual" => class_labels,
18     [predicted_classes[i] => round.(sample_preds[:, i], digits=2) for i in
19 1:10]...;
20     makeunique=true
21   )
22
23   PlutoUI.HTML("""
24   <div style="overflow-x: auto;">
25     <table>
26       <thead>
27         <tr>
28           $(join(["<th>$(col)</th>" for col in names(df)], " "))
29         </tr>
30       </thead>
31       <tbody>
32         $(join(["<tr>" * join(["<td>$(df[row, col])</td>" for col in
33 names(df)], " ") * "</tr>" for row in 1:size(df, 1)], "\n"))
34       </tbody>
35     </table>
36   </div>
37   """)
38 end
39
```

## Tip

Here's some of the code needed to create the DataFrame:

```
DataFrame(round.(model(rand_test), digits=2),  
          Symbol.(rand_label),  
          makeunique=true)
```

```
actual_classes =
```

```
["automobile", "deer", "cat", "automobile", "truck", "truck", "airplane", "frog", "airp
```

```
1 actual_classes = [class_labels[rand_labels[i] + 1] for i in 1:10]
```



```
1 @bind inx Slider(1:10, default=1)
```

Predicted: truck  
Actual: airplane







```
1 begin
2     selected_image = rand_test[:, :, :, inx]
3     selected_label_index = rand_labels[inx]
4     selected_label = class_labels[selected_label_index + 1]
5
6     selected_pred = sample_preds[:, inx]
7     selected_pred_index = onecold(selected_pred, 0:9)
8     selected_pred_class = class_labels[selected_pred_index + 1]
9
10    display_image = colorview(RGB, permutedims(cpu(selected_image), (3, 1, 2)))
11    plt = plot(display_image, axis=false, title="Predicted:
$selected_pred_class\nActual: $selected_label")
12
13    # display(plt)
14 end
```

**Model a:**





Classes/ Actual	automobile	deer	cat	automobile_1	truck	truck_1	airplane	frog	air
airplane	0.0	0.0	0.0	0.0	0.0	0.0	0.56	0.0	0.0
automobile	1.0	0.0	0.0	0.99	0.0	0.0	0.3	0.0	0.0
bird	0.0	0.01	0.31	0.0	0.0	0.0	0.0	0.0	0.0
cat	0.0	0.01	0.51	0.0	0.0	0.0	0.0	0.09	0.0
deer	0.0	0.71	0.03	0.0	0.0	0.0	0.0	0.01	0.0
dog	0.0	0.01	0.04	0.0	0.0	0.0	0.0	0.02	0.0
frog	0.0	0.0	0.1	0.0	0.0	0.0	0.0	0.88	0.0
horse	0.0	0.27	0.01	0.0	0.0	0.0	0.0	0.0	0.0
ship	0.0	0.0	0.0	0.0	0.0	0.0	0.14	0.0	0.0
truck	0.0	0.0	0.0	0.01	0.99	1.0	0.0	0.0	0.0

## Model b:

Classes/ Actual	automobile	deer	cat	automobile_1	truck	truck_1	automobile_2	dog
airplane	0.0	0.01	0.01	0.0	0.01	0.0	0.05	0.0
automobile	0.98	0.0	0.03	1.0	0.01	0.02	0.84	0.0
bird	0.0	0.02	0.06	0.0	0.0	0.0	0.0	0.0
cat	0.0	0.01	0.51	0.0	0.0	0.0	0.0	0.25
deer	0.0	0.74	0.05	0.0	0.0	0.0	0.0	0.03
dog	0.0	0.02	0.33	0.0	0.0	0.0	0.0	0.45
frog	0.0	0.0	0.02	0.0	0.0	0.0	0.0	0.15
horse	0.0	0.18	0.02	0.0	0.0	0.0	0.0	0.03
ship	0.0	0.0	0.0	0.0	0.0	0.0	0.03	0.0
truck	0.02	0.0	0.0	0.0	0.98	0.98	0.08	0.0

This looks similar to how we would expect the results to be. At this point, it's a good idea to see how our net actually performs on new data, that we have prepared.

## Overall accuracy

We iterate over the entire test set to calculate the overall model accuracy.

### Baseline Model:







```
1 begin
2     class_correct = zeros(10)
3     class_total = zeros(10)
4
5     for i in 1:length(test)
6         # Get predictions and labels for the batch
7         preds = model(test[i][1]) # preds: (10, batch_size) on GPU
8         lab = test[i][2] # lab: OneHotMatrix on GPU
9
10        # Move predictions to CPU and get predicted class indices
11        pred_classes = onecold(cpu(preds), 0:9) # Convert to class indices
12
13        actual_classes = Int.(cpu(lab.indices)) .- 1 # Convert UInt32 indices
14        to Int and adjust for 0-based indexing
15
16        # Update counts
17        for c in 0:9
18            idx = (actual_classes .== c)
19            class_total[c + 1] += sum(idx)
20            class_correct[c + 1] += sum(pred_classes[idx] .== c)
21        end
22    end
23 end
```







	accuracy	class
1	0.623	"airplane"
2	0.822	"automobile"
3	0.497	"bird"
4	0.359	"cat"
5	0.616	"deer"
6	0.571	"dog"
7	0.718	"frog"
8	0.768	"horse"
9	0.731	"ship"
10	0.673	"truck"

```
1 DataFrame(accuracy=(class_correct ./ class_total), class=classes)
```

0.615

```
1 round(mean([accuracy(test[i]...) for i in 1:10]), digits=3)
```

**Model a:**











```
1 begin
2   class_correct_a = zeros(10)
3   class_total_a = zeros(10)
4
5   for i in 1:length(test)
6     preds_a = model_a(test[i][1])
7     lab = test[i][2]
8
9     pred_classes_a = onecold(cpu(preds_a), 0:9)
10
11    actual_classes = Int.(cpu(lab.indices)) .- 1
12
13    for c in 0:9
14      idx = (actual_classes .== c)
15      class_total_a[c + 1] += sum(idx)
16      class_correct_a[c + 1] += sum(pred_classes_a[idx] .== c)
17    end
18  end
19 end
20
```

df\_model\_a =

	accuracy	class
1	0.661	"airplane"
2	0.706	"automobile"
3	0.529	"bird"
4	0.459	"cat"
5	0.711	"deer"
6	0.615	"dog"
7	0.642	"frog"
8	0.705	"horse"
9	0.783	"ship"
10	0.787	"truck"

```
1 df_model_a = DataFrame(accuracy=(class_correct_a ./ class_total_a),
  class=classes)
```

mean\_accuracy\_a = 0.66

## Model b:









```
1 begin
2     class_correct_b = zeros(10)
3     class_total_b = zeros(10)
4
5     for i in 1:length(test)
6         preds_b = model_b(test[i][1])
7         lab = test[i][2]
8
9         pred_classes_b = onecold(cpu(preds_b), 0:9)
10
11         actual_classes = Int.(cpu(lab.indices)) .- 1
12
13         for c in 0:9
14             idx = (actual_classes .== c)
15             class_total_b[c + 1] += sum(idx)
16             class_correct_b[c + 1] += sum(pred_classes_b[idx] .== c)
17         end
18     end
19 end
```













df\_model\_b =

	accuracy	class
1	0.661	"airplane"
2	0.685	"automobile"
3	0.496	"bird"
4	0.455	"cat"
5	0.632	"deer"
6	0.601	"dog"
7	0.704	"frog"
8	0.689	"horse"
9	0.619	"ship"
10	0.833	"truck"

---

10 0.000 track













```
1 df_model_b = DataFrame(accuracy=(class_correct_b ./ class_total_b),  
  class=classes)
```

```
mean_accuracy_b = 0.638
```















